

Implementation of Materialized Views in Cassandra

Alex Correia-Santos

Keywords:

Materialized Views Cassandra Key-Value Stores

1. Types of Views

1.1. Join View

```
CREATE JOIN <AB> OF <CF_A>.<COL> AND <CF_B>
```

Creates a Column Family (named <AB>) which stores the keys of the column family <CF_A> as columns for the key <COL>. Everytime a insertion is made into the base table <CF_A>, or a <COL> is updated, the view will be automatically updated to match the changes.

1.2. Index View

```
CREATE INDEX VIEW <A_v> OF <CF_A>.<COL_v>
```

Creates a copy of the column family <CF_A> with the values of the column <COL_v> acting as the key instead. As <COL_v> might have repeated values, the columns will be stored in a set-like fashion.

1.3. Non-key Join View

Such view can be obtained by creating an `Index View` then a `Join View`.

2. Views Properties

The following properties face different advantages and drawbacks, therefore they should be tested and changed in order to get the most efficient solutions.

- **Read-only**

Views cannot be changed directly by the user, the only functions that can be executed over a view are **GET** and **DROP**. (**pros:** guaranteed integrity of the view, **cons:** extra “check” to every function)

- **Data Duplication**

All the columns inserted into the base table will be automatically inserted into the related views as **index:column_name**. (**pros:** faster information retrieval, **cons:** slower data insertion)

- **Cascade Drop**

Whenever a base table is dropped, every view related to such base table shall be dropped aswell (**pros:** less space taken by (what might be) useless data, **cons:** slower drops)

3. Consistency

- **Insert into base view ->Insert into materialized view**

Guarantees that materialized view will only be updated if base view is successfully updated. In case of idempotent-only operations, it's guaranteed eventual consistency, as updates on materialized view will be constantly pushed until it goes through, or until a newer update is committed.

4. Consistency Issues Solutions

- **Full Lock**

Only one write a time, guarantees a full consistency between views but will decrease the performance drastically and affect directly the goal of a key-value stores.

- **Lock field**

If possible, locking only a field can guarantee that both views are in the same state before allowing any other operation to be made on that field. It might decrease the database's performance.

- **Test** *timestamp* \rightarrow *Try* **update** (currently on use)

Repeatdly try to update materialized views until the update is successful or a higher timestamp update is committed.